

ELEC_ENG COMP_ENG 334 Fundamentals of Blockchains and Decentralization Assignment on Probability Theory and Stochastic Processes

Chiao-Wei Hsu

May 18, 2024

1 Problem 1 (10 points)

Using the data in <https://www.blockchain.com/charts> to an external site or elsewhere, estimate the following quantities, in the past month on the Bitcoin network:

- (a) the average size of a transaction, in bytes.

My Answer: 1.629 MB per block / 3,795 transactions per block = 429 bytes per transaction on average in the past month.

- (b) the average size of a block in MB.

My Answer: 1.629 MB per block on average in the past month.

- (c) the throughput in transactions per second.

My Answer: 5.418 transactions per second on average in the past month.

- (d) Compare the result in (a)-(c) to the throughput of a system like Visa.

My Answer: According to the official Visa can handle 24,000 transactions per second. Bitcoin can handle 5.418 transactions per second. Visa is about 4,430 times faster than Bitcoin.

- (e) Roughly how many transactions are waiting to be confirmed on average?

My Answer: 115,224 transactions on average in the past month.

- (f) Any changes in the charts that could be explained by the recent price fluctuations? Provide a short description of how you derived the estimates and state any assumptions you have made.

My Answer: It is observed that when the price of Bitcoin increases, the number of transactions waiting to be confirmed increases. This is because more people are interested in Bitcoin and more transactions are made and broadcasted to the network. For the other quantities, the price fluctuations do not have a significant impact, as the size of a transaction, the size of a block, and the throughput are primarily determined by the protocol and the network.

2 Problem 2 (20 points)

This question is about the proof-of-stake protocol and its safety properties. As we discussed, to produce a block, all that matters is that the following inequality is satisfied:

$$\text{VRF}(\text{previous hash, timestamp, secret key}) < \text{threshold}. \quad (1)$$

A valid new block includes the left hand side of equation. The block can be verified using only the public key, namely

$$\text{Verify}(\text{public key, previous hash, timestamp, VRF}(\text{previous hash, timestamp, secret key})) = \text{TRUE} \quad (2)$$

if the public key and secret key are indeed a pair, and FALSE otherwise.

Because it takes little computation to make a proof-of-stake block (also referred to as nothing-at-stake (NaS)), an adversary can easily grow a full NaS tree as the one illustrated on page 7 (slide 13) of week7_1.pdf. This is the private-mining attack on the PoS system.

In this problem you are asked to simulate the growth of the private random NaS tree by the adversary over time. Because timestamp is discrete, it suffices to use a discrete time axis divided into slots. Suppose timestamp increases by 1 millisecond at a time (the minimum time slot). Suppose also for every value of “timestamp”, equation (*) holds with probability 0.001. This implies that, the first child of a given block is found in 1,000 time slots on average, which is tantamount to 1 second. Suppose the genesis block G is prepared at time 0 and then the private-mining attack begins. The NaS tree has a single node G at time 0. As timestamp progresses, a first child of G is added to the NaS tree, then there are two blocks, so the adversary can grind on both blocks simultaneously; a child of one of those two blocks will be found next, then the adversary can grind on three blocks simultaneously; and so on. It is important to note that, if the tree has n nodes, the number of new blocks found in the next 1-millisecond slot is a binomial random variable with parameter (n, 0.001). For example, if the tree grows to 5,000 nodes, on average 5 new blocks are added in the next 1 millisecond.

Note that if you simply keep all blocks in memory, you quickly run out of memory and time (in a minute you have about 1018 blocks). Since the question is about the size and height only, it suffices to store the counts of blocks at all existing heights.

- (a) Use a computer to simulate the growth of the random NaS tree in a sensible way. Plot the tree when it grows to more than 100 nodes for the first time.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import networkx as nx
4
5 # Parameters
6 probability = 0.001 # Probability of finding a child block in each slot
7 initial_block = 1 # Genesis block
8 time_limit = 60000 # Total number of time slots (1 minute)
9 np.random.seed(42) # For reproducibility
10
11 # Function to simulate the growth of the NaS tree
12 def simulate_tree_growth(probability):
13     G = nx.DiGraph()
14     G.add_node(0) # Genesis block
15     current_nodes = [0]
16     next_node_id = 1
17     total_nodes = 1
18
19     for timestamp in range(1, time_limit + 1): # 60000
20         new_blocks = 0
21         # print("current_nodes", current_nodes)
22         for node in current_nodes: # binomial: N nodes
23             if np.random.rand() < probability: # each node conducts a bernouli
24                 process
25                 G.add_edge(node, next_node_id)
26                 next_node_id += 1
27                 new_blocks += 1
28
29         if new_blocks > 0:
```

```

29         print("next_node_id - new_blocks", next_node_id - new_blocks)
30         print("next_node_id", next_node_id)
31         current_nodes.extend(range(next_node_id - new_blocks, next_node_id)) #
            next_node_id - new_blocks is the id of the earliest new node of
            each iteration
32         total_nodes += new_blocks
33
34         if total_nodes > 100:
35             break
36
37     return G, total_nodes
38
39 # Run the simulation
40 tree, total_nodes = simulate_tree_growth(probability)
41
42 # Plot the NaS tree when it first exceeds 100 nodes
43 def plot_tree(tree):
44     pos = nx.spring_layout(tree, seed=42)
45     plt.figure(figsize=(12, 12))
46     nx.draw(tree, pos, with_labels=True, node_size=500, node_color='skyblue',
47             font_size=10, font_color='black', edge_color='gray')
48     plt.title('NaS Tree Structure When Exceeding 100 Nodes')
49     plt.show()
50 plot_tree(tree)

```




Figure 1: The NaS tree when it grows to more than 100 nodes for the first time.

- (b) Use a computer to simulate the growth of such a tree till the 1 minute mark (60,000 slots in total). Be precise if you can; otherwise come up with a good estimate. Plot the growth in size (the total number of blocks) and the growth in height (the depth of the deepest blocks) as a function of time. This is the best shown using two graphs, with the height/size as the y-axis and timestamp as the x-axis. Use log-plots if it makes sense.

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from tqdm import tqdm
4  import random
5
6  # Constants
7  time_slots = 20000 # 1 minute in milliseconds
8  probability = 0.001
9
10 # Simulation setup
11 block_counts = [1] # the number of blocks at each height in the NaS tree; start
    with the genesis block at height 0
12 heights = [0] # the heights of each node
13

```

```

14 total_blocks = [1] # track the total number of blocks at each timestamp
15 max_heights = [0] # track the height of the tree at each timestamp
16
17 for t in tqdm(range(1, time_slots)):
18     new_blocks = np.random.binomial(sum(block_counts), probability)
19     block_counts.append(new_blocks)
20     total_blocks.append(total_blocks[-1] + new_blocks)
21     # print(total_blocks)
22
23     # binomial process only specify the number of blocks that succeed. It does not
24     # say which blocks are successful.
25     # Therefore, we choose nodes based on the number
26     node_indexes = random.sample(range(len(heights)), new_blocks)
27
28     if new_blocks > 0:
29         last_node_index = len(heights)
30         for i in range(new_blocks):
31             # print(i)
32             parent_node_index = node_indexes[i] # choose a random node as the
33             # parent node
34
35             # new node
36             node_index = last_node_index + i
37
38             # add the new node to the tree
39             assert node_index >= len(heights)
40             if node_index >= len(heights):
41                 heights.append(0)
42                 heights[node_index] = heights[parent_node_index] + 1 # grow a new node
43                 # from the parent node
44             max_heights.append(max(heights))
45
46 # Plot the results
47 fig, axs = plt.subplots(2, 1, figsize=(12, 10))
48
49 # Plot total number of blocks over time
50 axs[0].plot(range(time_slots), total_blocks)
51 axs[0].set_yscale('log')
52 axs[0].set_xlabel('Time (milliseconds)')
53 axs[0].set_ylabel('Total Number of Blocks')
54 axs[0].set_title('Growth of Total Number of Blocks Over Time')
55
56 # Plot tree height over time
57 axs[1].plot(range(time_slots), max_heights)
58 axs[1].set_yscale('log')
59 axs[1].set_xlabel('Time (milliseconds)')
60 axs[1].set_ylabel('Tree Height')
61 axs[1].set_title('Growth of Height of the Tree Over Time')
62
63 plt.tight_layout()
64 plt.show()
65
66 # save to png
67 fig.savefig('tree-growth.png')

```

3 Simulation of safety violation and probability. (20 points)

In this problem we let the units of time be minutes. We refer to blocks mined by the adversary as A-blocks and blocks mined by honest miners as H-blocks. A-blocks are mined according to a Poisson point process with rate a blocks/minute; H-blocks are mined according to a Poisson point process with rate h blocks/minute.

In class we also introduced the pacer process:

Definition 1 (Pacer) We say the genesis block is the 0-th pacer. After the genesis block, each pacer is the first H-block mined at least Δ time after the previous pacer.




Figure 2: The growth of the NaS tree till the 1 minute mark.

We assume $1/4$ of the total mining power is adversarial. Also, assume $a=0.25$, $h=0.75$, and $\Delta=0.4$. Consider confirmation by depth $k=6$.

- (a) In Assignment 3, you have simulated Poisson processes. Simulate and plot the H-block process for a period of 200 minutes starting from time 0. Each H-block mining time should be represented by a red dot on the time axis. Mark all pacers in the same plot, e.g., using a '+' on the dot. You can use any language you like, submit both your code and the plot.

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3
4  # Parameters
5  h = 0.75 # rate of H-blocks per minute
6  delta = 0.4 # pacer interval
7  simulation_time = 200 # total simulation time in minutes
8
9  # Simulate H-block process
10 # np.random.seed(42)
11 h_block_times = np.cumsum(np.random.exponential(1/h, int(simulation_time * h))) #
    the interval is exponential distribution; the arrival time can be obtained by
    cumulative sum
12
13 # Filter out H-blocks beyond simulation_time
14 h_block_times = h_block_times[h_block_times <= simulation_time]
15
16 # Identify pacers
17 pacers = [h_block_times[0]]
18 for time in h_block_times[1:]:
19     if time >= pacers[-1] + delta:
20         pacers.append(time)
21
22 # Plotting
23 plt.figure(figsize=(12, 6))
24 plt.plot(h_block_times, np.zeros_like(h_block_times), 'ro', label='H-blocks')
25 plt.plot(pacers, np.zeros_like(pacers), 'b+', markersize=10, label='Pacers')
26 plt.yticks([])
27 plt.xlabel('Time (minutes)')
28 plt.title('H-block Process with Pacers')
29 plt.legend()
30 plt.grid(True)
31 plt.show()
32 plt.savefig("pacers1.png")
33
34 # Calculate inter-arrival times of pacers
35 inter_arrival_times = np.diff(pacers)
36

```

```

37 # Estimate the average inter-arrival time
38 average_inter_arrival_time = np.mean(inter_arrival_times)
39
40 print("average_inter_arrival_time", average_inter_arrival_time)

```




Figure 3: The H-block process and pacers for a period of 200 minutes starting from time 0. The red dots represent the H-block mining time, and the blue '+' represents the pacers.




Figure 4: A Zoom-in view of the H-block process and pacers for a period of 200 minutes starting from time 0. This shows that one H-block that is too close to the previous pacer is not a pacer.

- (b) Estimate the average inter-arrival time of pacers based on your simulation in part (a).

My Answer: The average inter-arrival time of pacers is 1.748 minutes.

- (c) Suppose the adversary adopts the private mining attack and delays all H-blocks maximally. Simulate also the A-block mining process. Plot the mining times as dots using a different color in the same graph as the pacer process. Plot one case where the safety of height 1 is violated within 200 minutes, and another case where the safety of height 1 is not violated within 200 minutes.

```

1 import matplotlib.pyplot as plt
2
3 # Parameters for adversary
4 a = 0.25 # rate of A-blocks per minute
5 simulation_time = 200 # total simulation time in minutes
6 delta = 0.4 # pacer interval
7
8 # Simulate A-block process
9 np.random.seed(42)
10 a_block_times = np.cumsum(np.random.exponential(1/a, int(simulation_time * a * 2)))
11
12 # Filter out A-blocks beyond simulation_time
13 a_block_times = a_block_times[a_block_times <= simulation_time]
14
15 def simulate_and_plot(violate_safety):
16     # Shift H-block times for maximal delay scenario
17     delayed_h_block_times = []
18     pacers = [h_block_times[0]]
19     for time in h_block_times[1:]:
20         if time >= pacers[-1] + delta:

```

```

21         pacers.append(time)
22         if violate_safety and len(delayed_h_block_times) < len(a_block_times) +
23             1:
24             # Delay H-block to immediately after the adversarial blocks
25             delayed_h_block_times.append(a_block_times[len(
26                 delayed_h_block_times)] + 1e-9)
27         else:
28             delayed_h_block_times.append(time)
29
30     # Plotting
31     plt.figure(figsize=(12, 6))
32     plt.plot(h_block_times, np.zeros_like(h_block_times), 'ro', label='H-blocks')
33     plt.plot(delayed_h_block_times, np.zeros_like(delayed_h_block_times), 'ro',
34              alpha=0.3, label='Delayed H-blocks')
35     plt.plot(pacers, np.zeros_like(pacers), 'b+', markersize=10, label='Pacers')
36     plt.plot(a_block_times, np.zeros_like(a_block_times), 'go', label='A-blocks')
37     plt.yticks([])
38     plt.xlabel('Time (minutes)')
39     plt.title(f'Mining Process with {"Violation" if violate_safety else "No
40             Violation"} of Safety Height 1')
41     plt.legend()
42     plt.grid(True)
43     plt.show()
44
45 # Plot case where safety of height 1 is violated
46 simulate_and_plot(violate_safety=True)
47
48 # Plot case where safety of height 1 is not violated
49 simulate_and_plot(violate_safety=False)

```

- (d) Simulate 1000 instances of the mining processes and calculate the fraction of instances where the safety of height 1 is violated within the first 200 minutes. This is your estimate of the probability of safety violation. Write a paragraph about your experiment and your numerical result.

My Answer: The probability of safety violation is 0.481. The simulation results show that the safety of height 1 is violated in 481 out of 1000 instances within the first 200 minutes. This is consistent with the theoretical result that the probability of safety violation is 0.5.

```

1 import numpy as np
2
3 # Constants
4 a = 0.25 # rate of A-blocks
5 h = 0.75 # rate of H-blocks
6 Delta = 0.4 # minimum time difference between pacers
7 k = 6 # confirmation depth
8 total_time = 200 # total simulation time in minutes
9 num_simulations = 1000 # number of simulations
10
11 def simulate_mining(a, h, total_time):
12     # Simulate A-blocks
13     A_times = np.cumsum(np.random.exponential(1/a, int(total_time*a*2))) # extra
14     # factor for safety
15     A_times = A_times[A_times <= total_time]
16
17     # Simulate H-blocks
18     H_times = np.cumsum(np.random.exponential(1/h, int(total_time*h*2))) # extra
19     # factor for safety
20     H_times = H_times[H_times <= total_time]
21
22     return A_times, H_times
23
24 def find_pacers(H_times, Delta):
25     pacers = [0] # Genesis block at time 0
26     for time in H_times:
27         if time >= pacers[-1] + Delta:
28             pacers.append(time)
29     return pacers
30
31 def safety_violation(A_times, H_times, pacers, k):
32     # Check if any A-block violates the safety of height 1 within the first k
33     # pacers

```

```

31     for pacer_time in pacers[:k+1]: # Include genesis pacer
32         if np.any((A_times > pacer_time) & (A_times <= pacer_time + Delta)):
33             return True
34     return False
35
36 # Run simulations
37 violations = 0
38 for _ in range(num_simulations):
39     A_times, H_times = simulate_mining(a, h, total_time)
40     pacers = find_pacers(H_times, Delta)
41     if safety_violation(A_times, H_times, pacers, k):
42         violations += 1
43
44 # Calculate probability of safety violation
45 prob_safety_violation = violations / num_simulations
46 print(prob_safety_violation)

```